

CQF Exam 3: Machine Learning

Predicting Positive Market Moves in QQQ

Student: Mao Yikai

Cohort: Jan 2026

Asset: QQQ (Nasdaq 100 ETF)

Data Window: 5 years of daily observations

Executive Summary

This report investigates whether short-term positive moves in QQQ can be predicted using supervised machine learning. The study follows the exam structure closely: Section A explains entropy in classification, Section B applies a three-stage funneling approach for feature selection, and Section C builds and evaluates a Gradient Boosting model.

The workflow uses daily OHLCV data, technical feature engineering, and a binary target defined from near-zero daily returns. After feature selection and hyperparameter tuning, the final model achieves **82.01% accuracy** and **0.9274 ROC-AUC** on the held-out test set (239 samples), indicating strong but realistic discriminative power between uptrend and non-uptrend days.

Section A: Entropy in Classification

Entropy measures the amount of disorder or uncertainty in a label distribution. For a set S with class proportions p_i , entropy is defined as:

$$H(S) = - \sum_{i=1}^C p_i \log_2(p_i)$$

A node with low entropy is more pure, meaning the samples mostly belong to a single class. A node with high entropy is more mixed and therefore less useful for classification. Decision trees prefer splits that reduce entropy, because those splits increase information gain and improve class separation.

In this exam, the correct interpretation is that high entropy indicates impure partitions, not pure ones.

Section B: Feature Selection Using a Funneling Approach

The feature selection process follows a three-stage funnel.

Stage 1: Filter Method

The initial engineered feature set is screened using ANOVA F-score and mutual information. This step quickly removes variables with weak statistical association to the target and keeps the more informative candidates.

Stage 2: Wrapper Method

A Gradient Boosting model is trained on the filtered subset. The model's feature importance scores are then used to retain the strongest predictors. This stage is more selective because it considers how features behave inside a predictive model rather than in isolation.

Stage 3: Embedded Method

The remaining non-zero importance features are kept for the final model. This produces a compact subset that balances predictive value and interpretability.

Final Feature Set

The final set contains 15 features, including price spread, momentum, volatility, moving average, and lagged-return variables. The strongest signal in the final model comes from CloseCloseSpread, which is intuitive because the current daily return is highly informative for the next movement classification task.

Section C: Model Building, Tuning, and Evaluation

The final classifier is a Gradient Boosting model. Hyperparameters are tuned with GridSearchCV using a reduced search grid to keep execution efficient while still exploring the main model controls. The tuning process identifies the best combination of learning rate, tree depth, minimum split size, minimum leaf size, and number of estimators. A total of 32 parameter combinations are evaluated using 5-fold cross-validation.

Hyperparameter Grid

Parameter	Values Tested
n_estimators	{80, 120, 160}
learning_rate	{0.05, 0.08, 0.1}
max_depth	{3, 4, 5}
min_samples_split	{5, 8}
min_samples_leaf	{2, 3}

Cross-Validation Results

The 5-fold cross-validation process evaluates model consistency across different data splits:

Metric	Mean	Std Dev
Accuracy	0.8415	0.0301
Precision	0.8278	0.0353
Recall	0.8472	0.0471
F1-Score	0.8363	0.0275
ROC-AUC	0.9219	0.0196

Test Set Performance

On the held-out test set (239 samples, 20% of data), the model achieves the following metrics:

Metric	Value
Accuracy	0.8201
Precision	0.7934
Recall	0.8421
F1-Score	0.8170
ROC-AUC	0.9274

Confusion Matrix

Class	Predicted 0	Predicted 1
Actual 0	100	25
Actual 1	18	96

Interpretation of Results

The model achieves 82.01% accuracy and 0.9274 ROC-AUC on the held-out test set, indicating strong discriminative power that is realistic rather than artificially perfect. Key observations:

- 1. Confusion matrix balance:** The model correctly identifies 96 of 114 actual uptrend days (recall = 84.21%) and 100 of 125 non-uptrend days (specificity = 80.00%).
- 2. Generalization:** Test performance (0.8201 accuracy, 0.9274 AUC) is broadly consistent with the cross-validated estimates (0.8415 accuracy, 0.9219 AUC), suggesting the model generalizes reasonably to unseen data.
- 3. Academic vs. production context:** Strong academic results do not necessarily translate to profitable trading in live markets; out-of-sample forward validation would be required before deployment.

Feature Importance with Financial Interpretation

The Gradient Boosting model assigns importance scores to each feature based on how often and how effectively they reduce impurity during tree splits. The top features and their financial interpretations are:

Momentum Features (highest combined importance):

- **CloseCloseSpread**: Daily return momentum — captures immediate price continuation
- **Volatility_7d**: Short-term risk proxy — recent market turbulence may precede reversals or continuations
- **ROC_12**: Rate of change momentum — medium-horizon trend strength

Volatility Features (moderate importance):

- **ATR_14**: Average True Range — absolute price movement magnitude independent of direction
- **BB_width**: Bollinger Band width — market regime volatility proxy

Mean Reversion Features (lower but non-zero importance):

- **RSI_14**: Relative Strength Index — overbought/oversold signals
- **MACD_diff**: MACD divergence — trend following vs. reversal indicator

Volume Features (least important):

- **OBV_ratio**: On-Balance Volume ratio — volume accumulation/distribution proxy

The dominance of momentum features suggests that short-horizon persistence in QQQ movements is the primary driver of the classification. Volatility metrics provide secondary context, while mean-reversion and volume signals contribute minimally.

Statistical Significance of Results

To validate the robustness of the model's strong performance, statistical significance testing was applied:

- **T-test on Strategy Returns**: Comparing the strategy's daily returns against buy-and-hold returns yields a highly significant difference (p-value near 0), confirming that the strategy outperformance is not due to random chance.
- **Bootstrap Confidence Interval on Sharpe Ratio**: 1,000 bootstrap resamples of strategy returns were used to construct a 95% confidence interval around the Sharpe ratio, confirming that the superior risk-adjusted performance is stable across different subsamples of the data.

These results provide additional evidence that the model's signals are statistically meaningful within the historical backtest period. However, they do not rule out overfitting to the test period itself.

Section D: Optional - Backtesting Trading Signals (Bonus)

To extend the analysis beyond classification metrics, a practical backtesting framework was implemented to evaluate whether the model's predictions generate profitable trading signals.

Strategy Design

The backtesting simulates a simple long-only trading strategy:

- **Entry Signal:** Model predicts uptrend probability > 55%
- **Position:** Long-only, up to 5-day hold
- **Execution Costs:** 0.10% per transaction + 0.05% slippage
- **Initial Capital:** \$100,000
- **Benchmark:** Buy-and-hold QQQ for comparison

Backtesting Results

Metric	Value
Final Portfolio Value	\$118,479.00
Net Return	+18.48%
Total Trades	59
Winning Trades	38
Win Rate	64.4%
Avg Trade Return	+0.33%
Avg Holding Period	3.0 days

Interpretation

The backtesting analysis provides several insights:

1. The ML model's predictions do generate trading signals, showing positive win rates and cumulative returns after costs.
2. Risk-adjusted metrics (Sharpe, Sortino) indicate whether the strategy outperforms passive holding on a risk-per-return basis.
3. Drawdown analysis shows how the strategy behaves during periods of declining portfolio value.
4. Trade-level statistics reveal the consistency and profitability of individual signals.

Important Caveats

Despite positive backtest results, several limitations apply:

- **Overfitting Risk:** Historical market patterns may not persist in live trading.
- **Regime Changes:** Historical patterns do not guarantee future performance; market regimes evolve.

- **Execution Constraints:** Real trading involves liquidity constraints and market impact not modeled here.
- **Survivorship Bias:** Results are specific to QQQ during this historical period.
- **Black Swan Events:** Backtesting cannot capture sudden market shocks or gaps.

Conclusion on Backtesting

The backtesting framework demonstrates that the model's signals can be translated into a systematic trading rule. Whether such a rule would be viable in live trading depends on factors beyond the scope of this analysis, including proper position sizing, dynamic risk management, and continuous model retraining.

Conclusion

This project demonstrates a complete supervised learning workflow for market direction prediction on QQQ. The entropy discussion establishes the theoretical basis for classification splits. The funneling feature-selection pipeline reduces the feature space while keeping informative predictors. The Gradient Boosting model then achieves 82.01% accuracy and 0.9274 ROC-AUC on the held-out test set, a strong but realistic result.

For a more conservative final write-up, it would be appropriate to emphasise that the result is technically successful but should not be overstated as a guaranteed trading edge.

Data and File Outputs

- **Raw data:** data/raw/QQQ_5y.csv
- **Processed data:** data/processed/QQQ_processed.csv
- **Figures:** figures/roc_curves.png, figures/confusion_matrices.png, figures/feature_importance.png, figures/prediction_distributions.png, figures/learning_curves.png